



```
In [1]: 1 from keras.datasets import mnist  
2 from keras import models  
3 from keras import layers  
4 (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
D:\Anaconda3\lib\site-packages\h5py\_init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated.  
In future, it will be treated as `np.float64 == np.dtype(float).type`.  
from ._conv import register_converters as _register_converters  
Using TensorFlow backend.
```

```
In [2]: 1 train_images = train_images.reshape((60000, 28 * 28))  
2 train_images = train_images.astype('float32') / 255  
3  
4 test_images = test_images.reshape((10000, 28 * 28))  
5 test_images = test_images.astype('float32') / 255  
6  
7 from keras.utils import to_categorical  
8  
9 train_labels = to_categorical(train_labels)  
10 test_labels = to_categorical(test_labels)
```

```
In [3]: 1 import matplotlib.pyplot as plt  
2 %matplotlib inline
```

```
In [4]: 1 #優化器rmsprop  
2 network = models.Sequential()  
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))  
4 network.add(layers.Dense(10, activation='softmax'))  
5 network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])  
6 rmsprop_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)  
7 rmsprop_test_loss, rmsprop_test_acc = network.evaluate(test_images, test_labels)  
8 print('test_loss:', rmsprop_test_loss)  
9 print('test_acc:', rmsprop_test_acc)  
10 print(rmsprop_hist.history)  
11  
12 plt.figure(figsize=(10,3))  
13 plt.suptitle('RMSprop', fontsize=18)  
14  
15 plt.subplot(1,2,1)  
16 plt.plot(rmsprop_hist.history['acc'], 'b', linewidth=3)  
17 plt.title('model accuracy')  
18 plt.xlabel('epoch')  
19 plt.ylabel('accuracy')  
20 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])  
21 plt.grid(True)  
22 plt.tight_layout()  
23  
24 plt.subplot(1,2,2)  
25 plt.plot(rmsprop_hist.history['loss'], 'g', linewidth=3)  
26 plt.title('model loss')  
27 plt.xlabel('epoch')  
28 plt.ylabel('loss')  
29 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])  
30 plt.grid(True)  
31 plt.tight_layout()  
32  
33 plt.savefig('Accuracy of RMSprop.jpg', dpi=300)
```

```
WARNING:tensorflow:From D:\Anaconda3\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
```

```
Instructions for updating:
```

```
Colocations handled automatically by placer.
```

```
WARNING:tensorflow:From D:\Anaconda3\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
```

```
Instructions for updating:
```

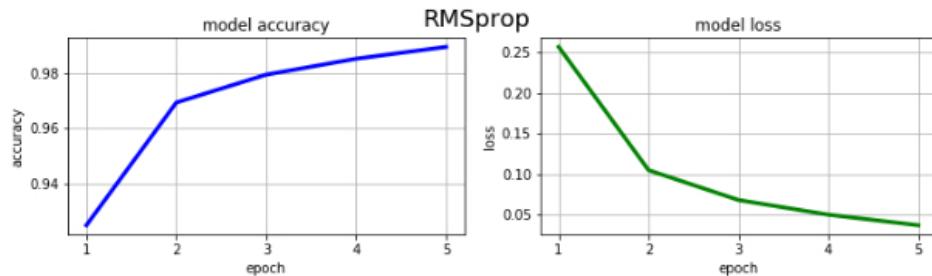
```
Use tf.cast instead.
```

```
Epoch 1/5
```

```

60000/60000 [=====] - 9s 152us/step - loss: 0.2565 - a
cc: 0.9248
Epoch 2/5
60000/60000 [=====] - 3s 52us/step - loss: 0.1048 - ac
c: 0.9692
Epoch 3/5
60000/60000 [=====] - 3s 53us/step - loss: 0.0681 - ac
c: 0.9793
Epoch 4/5
60000/60000 [=====] - 3s 51us/step - loss: 0.0500 - ac
c: 0.9851
Epoch 5/5
60000/60000 [=====] - 3s 55us/step - loss: 0.0371 - ac
c: 0.9894
10000/10000 [=====] - 1s 59us/step
test_loss: 0.06734216312763747
test_acc: 0.9799
{'loss': [0.2565254210035006, 0.10480482665896415, 0.06810588288803895, 0.05004
766948421796, 0.03708880497068167], 'acc': [0.9247833333333333, 0.9692333333015
442, 0.9793333333333333, 0.9851000000317891, 0.9893833333333333]}

```



```

In [5]: #優化器sgd
1 network = models.Sequential()
2 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
3 network.add(layers.Dense(10, activation='softmax'))
4 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['a
cc'])
5 sgd_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)
6 sgd_test_loss, sgd_test_acc = network.evaluate(test_images, test_labels)
7 print('test_loss:', sgd_test_loss)
8 print('test_acc:', sgd_test_acc)
9 print(sgd_hist.history)
10
11
12 plt.figure(figsize=(10,3))
13 plt.suptitle('SGD', fontsize=14)
14
15 plt.subplot(1,2,1)
16 plt.plot(sgd_hist.history['acc'], 'b', linewidth=3)
17 plt.title('model accuracy')
18 plt.xlabel('epoch')
19 plt.ylabel('accuracy')
20 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
21 plt.grid(True)
22 plt.tight_layout()
23
24 plt.subplot(1,2,2)
25 plt.plot(sgd_hist.history['loss'], 'g', linewidth=3)
26 plt.title('model loss')
27 plt.xlabel('epoch')
28 plt.ylabel('loss')
29 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
30 plt.grid(True)
31 plt.tight_layout()
32
33 plt.savefig('Accuracy of SGD.jpg',dpi=300)

```

```

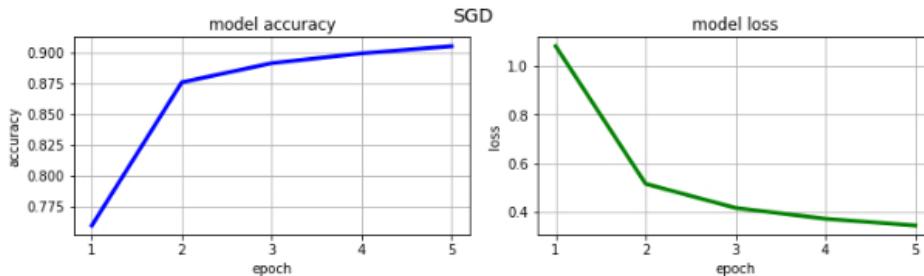
Epoch 1/5
60000/60000 [=====] - 3s 45us/step - loss: 1.0830 - ac
c: 0.7597
Epoch 2/5
60000/60000 [=====] - 3s 44us/step - loss: 0.5148 - ac
c: 0.8761
Epoch 3/5
60000/60000 [=====] - 3s 42us/step - loss: 0.4152 - ac
c: 0.8916
Epoch 4/5
60000/60000 [=====] - 3s 42us/step - loss: 0.3703 - ac
c: 0.8996
Epoch 5/5
60000/60000 [=====] - 3s 45us/step - loss: 0.3427 - ac

```

```

c: 0.9054
10000/10000 [=====] - 1s 62us/step
test_loss: 0.31555847700238226
test_acc: 0.9127
{'loss': [1.0829915905316672, 0.5147729955196381, 0.41524283391634625, 0.370302
22096443177, 0.3426549215475718], 'acc': [0.7596999999682108, 0.876050000031789
1, 0.891583333651225, 0.8995500000317892, 0.9053833333333333]}

```



In [6]:

```

1 #優化器adam
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 adam_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)
7 adam_test_loss, adam_test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', adam_test_loss)
9 print('test_acc:', adam_test_acc)
10 print(adam_hist.history)

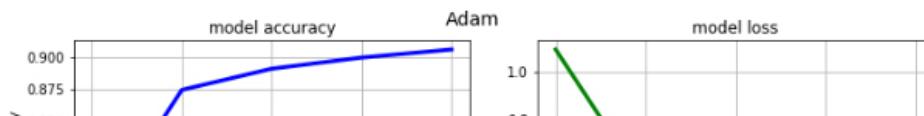
11
12 plt.figure(figsize=(10,3))
13 plt.suptitle('Adam', fontsize=14)
14
15 plt.subplot(1,2,1)
16 plt.plot(adam_hist.history['acc'], 'b', linewidth=3)
17 plt.title('model accuracy')
18 plt.xlabel('epoch')
19 plt.ylabel('accuracy')
20 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
21 plt.grid(True)
22 plt.tight_layout()
23
24 plt.subplot(1,2,2)
25 plt.plot(adam_hist.history['loss'], 'g', linewidth=3)
26 plt.title('model loss')
27 plt.xlabel('epoch')
28 plt.ylabel('loss')
29 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
30 plt.grid(True)
31 plt.tight_layout()
32
33 plt.savefig('Accuracy of Adam.jpg',dpi=300)

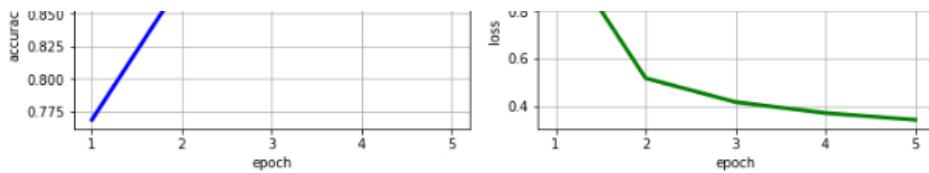
```

```

Epoch 1/5
60000/60000 [=====] - 4s 63us/step - loss: 1.0923 - accuracy: 0.7684
Epoch 2/5
60000/60000 [=====] - 3s 42us/step - loss: 0.5180 - accuracy: 0.8746
Epoch 3/5
60000/60000 [=====] - 3s 42us/step - loss: 0.4177 - accuracy: 0.8909
Epoch 4/5
60000/60000 [=====] - 3s 42us/step - loss: 0.3716 - accuracy: 0.8995
Epoch 5/5
60000/60000 [=====] - 3s 42us/step - loss: 0.3430 - accuracy: 0.9057
10000/10000 [=====] - 1s 61us/step
test_loss: 0.31735005323290827
test_acc: 0.9153
{'loss': [1.0923182662963866, 0.518003554871877, 0.4176862187862396, 0.3715899
117946625, 0.34304365968704226], 'acc': [0.7683500000317891, 0.874616666698455
8, 0.89095, 0.899533333015442, 0.9057499999682108]}

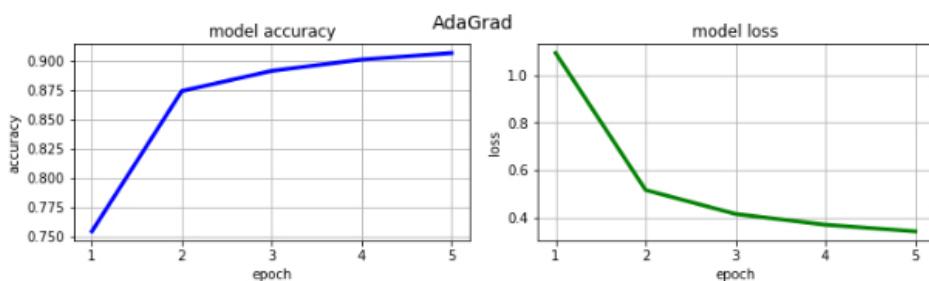
```





```
In [7]: 1 #優化器adagrad
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 adagrad_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)
7 adagrad_test_loss, adagrad_test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', adagrad_test_loss)
9 print('test_acc:', adagrad_test_acc)
10 print(adagrad_hist.history)
11
12 plt.figure(figsize=(10,3))
13 plt.suptitle('AdaGrad', fontsize=14)
14
15 plt.subplot(1,2,1)
16 plt.plot(adagrad_hist.history['acc'], 'b', linewidth=3)
17 plt.title('model accuracy')
18 plt.xlabel('epoch')
19 plt.ylabel('accuracy')
20 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
21 plt.grid(True)
22 plt.tight_layout()
23
24 plt.subplot(1,2,2)
25 plt.plot(adagrad_hist.history['loss'], 'g', linewidth=3)
26 plt.title('model loss')
27 plt.xlabel('epoch')
28 plt.ylabel('loss')
29 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
30 plt.grid(True)
31 plt.tight_layout()
32
33 plt.savefig('Accuracy of AdaGrad.jpg',dpi=300)
```

```
Epoch 1/5
60000/60000 [=====] - 3s 47us/step - loss: 1.0931 - acc: 0.7546
Epoch 2/5
60000/60000 [=====] - 3s 43us/step - loss: 0.5155 - acc: 0.8742
Epoch 3/5
60000/60000 [=====] - 3s 48us/step - loss: 0.4146 - acc: 0.8913
Epoch 4/5
60000/60000 [=====] - 3s 42us/step - loss: 0.3690 - acc: 0.9009
Epoch 5/5
60000/60000 [=====] - 3s 44us/step - loss: 0.3410 - acc: 0.9065
10000/10000 [=====] - 1s 61us/step
test_loss: 0.31547285937070846
test_acc: 0.9154
{'loss': [1.0930681081136067, 0.5154599870204926, 0.4146225591023763, 0.3690103939851125, 0.34104879151980083], 'acc': [0.7546166666348775, 0.87415, 0.8913166666666666, 0.9009333333333334, 0.906483333015442]}
```



```
In [8]: 1 #優化器adadelta
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```

6 adadelta_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)
7 adadelta_test_loss, adadelta_test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', adadelta_test_loss)
9 print('test_acc:', adadelta_test_acc)
10 print(adadelta_hist.history)

11 plt.figure(figsize=(10,3))
12 plt.suptitle('AdaDelta', fontsize=14)
13
14 plt.subplot(1,2,1)
15 plt.plot(adadelta_hist.history['acc'], 'b', linewidth=3)
16 plt.title('model accuracy')
17 plt.xlabel('epoch')
18 plt.ylabel('accuracy')
19 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
20 plt.grid(True)
21 plt.tight_layout()

22 plt.subplot(1,2,2)
23 plt.plot(adadelta_hist.history['loss'], 'g', linewidth=3)
24 plt.title('model loss')
25 plt.xlabel('epoch')
26 plt.ylabel('loss')
27 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
28 plt.grid(True)
29 plt.tight_layout()

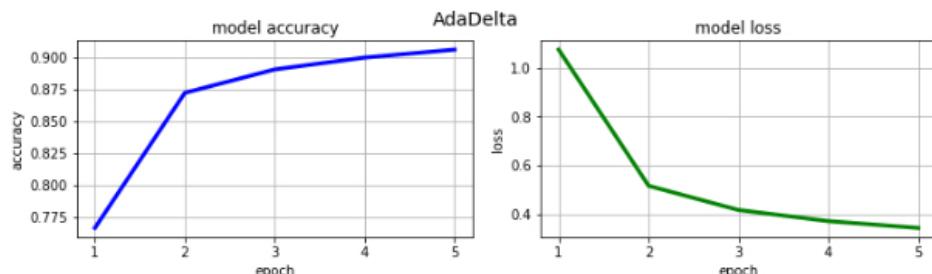
30 plt.savefig('Accuracy of AdaDelta.jpg', dpi=300)

```

```

Epoch 1/5
60000/60000 [=====] - 3s 49us/step - loss: 1.0738 - acc: 0.7666
Epoch 2/5
60000/60000 [=====] - 3s 44us/step - loss: 0.5171 - acc: 0.8720
Epoch 3/5
60000/60000 [=====] - 3s 43us/step - loss: 0.4181 - acc: 0.8902
Epoch 4/5
60000/60000 [=====] - 3s 43us/step - loss: 0.3727 - acc: 0.8995
Epoch 5/5
60000/60000 [=====] - 3s 43us/step - loss: 0.3448 - acc: 0.9057
10000/10000 [=====] - 1s 64us/step
test_loss: 0.31858390017151833
test_acc: 0.9103
{'loss': [1.0738442731539408, 0.5170751467227935, 0.4181419246514638, 0.37266617852846784, 0.34476230732599894], 'acc': [0.7665666666984559, 0.8719666666984558, 0.8902499999682109, 0.8994999999682108, 0.9057]}

```



```

In [9]: #優化器nadam
1 network = models.Sequential()
2 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
3 network.add(layers.Dense(10, activation='softmax'))
4 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
5 nadam_hist = network.fit(train_images, train_labels, epochs=5, batch_size=128)
6 nadam_test_loss, nadam_test_acc = network.evaluate(test_images, test_labels)
7 print('test_loss:', nadam_test_loss)
8 print('test_acc:', nadam_test_acc)
9 print(nadam_hist.history)

10 plt.figure(figsize=(10,3))
11 plt.suptitle('Nadam', fontsize=14)
12
13 plt.subplot(1,2,1)
14 plt.plot(nadam_hist.history['acc'], 'b', linewidth=3)
15 plt.title('model accuracy')
16 plt.xlabel('epoch')
17 plt.ylabel('accuracy')

```

```

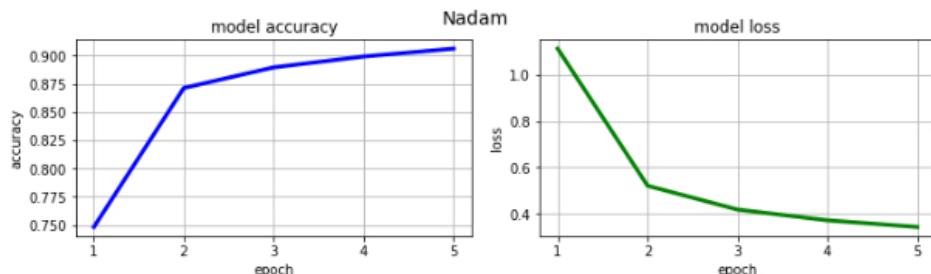
18 plt.xlabel('epoch')
19 plt.ylabel('accuracy')
20 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
21 plt.grid(True)
22 plt.tight_layout()
23
24 plt.subplot(1,2,2)
25 plt.plot(nadam_hist.history['loss'], 'g', linewidth=3)
26 plt.title('model loss')
27 plt.xlabel('epoch')
28 plt.ylabel('loss')
29 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
30 plt.grid(True)
31 plt.tight_layout()
32
33 plt.savefig('Accuracy of Nadam.jpg',dpi=300)

```

```

Epoch 1/5
60000/60000 [=====] - 3s 46us/step - loss: 1.1132 - ac
c: 0.7484
Epoch 2/5
60000/60000 [=====] - 3s 46us/step - loss: 0.5214 - ac
c: 0.8713
Epoch 3/5
60000/60000 [=====] - 3s 45us/step - loss: 0.4187 - ac
c: 0.8895
Epoch 4/5
60000/60000 [=====] - 2s 41us/step - loss: 0.3723 - ac
c: 0.8991
Epoch 5/5
60000/60000 [=====] - 3s 44us/step - loss: 0.3436 - ac
c: 0.9061
10000/10000 [=====] - 1s 63us/step
test_loss: 0.31700938050150873
test_acc: 0.913
{'loss': [1.113154947121938, 0.5213805438995361, 0.4187394819736481, 0.37233679
865201313, 0.3436127088069916], 'acc': [0.748383333651225, 0.871283333015442,
0.8895166666984559, 0.8990999999682109, 0.9060666666348776]}

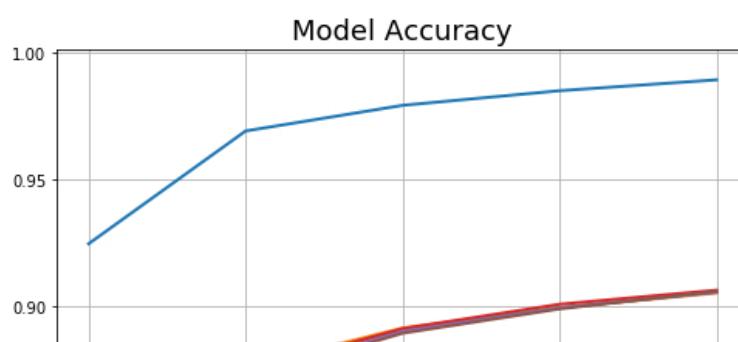
```

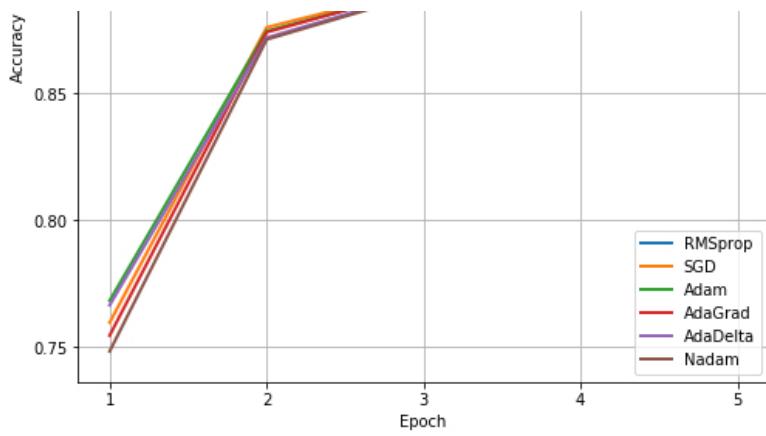


```

In [10]: 1 # summarize history for accuracy
2 plt.figure(figsize=(8,8))
3 plt.plot(rmsprop_hist.history['acc'], linewidth=2)
4 plt.plot(sgd_hist.history['acc'], linewidth=2)
5 plt.plot(adam_hist.history['acc'], linewidth=2)
6 plt.plot(adagrad_hist.history['acc'], linewidth=2)
7 plt.plot(adadelta_hist.history['acc'], linewidth=2)
8 plt.plot(nadam_hist.history['acc'], linewidth=2)
9 plt.title('Model Accuracy', fontsize=18)
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy')
12 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
13 plt.legend(['RMSprop', 'SGD', 'Adam', 'AdaGrad', 'AdaDelta', 'Nadam'], loc='t
14 plt.grid(True)
15 plt.savefig('Accuracy of different optimizer.jpg',dpi=300)

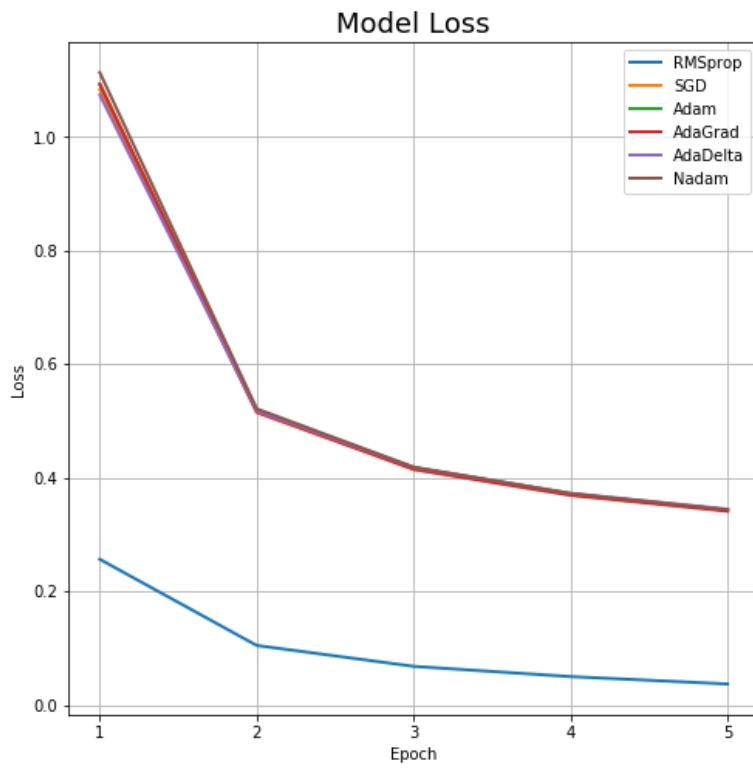
```





In [11]:

```
1 # summarize history for loss
2 plt.figure(figsize=(8,8))
3 plt.plot(rmsprop_hist.history['loss'], linewidth=2)
4 plt.plot(sgd_hist.history['loss'], linewidth=2)
5 plt.plot(adam_hist.history['loss'], linewidth=2)
6 plt.plot(adagrad_hist.history['loss'], linewidth=2)
7 plt.plot(adadelta_hist.history['loss'], linewidth=2)
8 plt.plot(nadam_hist.history['loss'], linewidth=2)
9 plt.title('Model Loss', fontsize=18)
10 plt.xlabel('Epoch')
11 plt.ylabel('Loss')
12 plt.xticks([0, 1, 2, 3, 4], ['1', '2', '3', '4', '5'])
13 plt.legend(['RMSprop', 'SGD', 'Adam', 'AdaGrad', 'AdaDelta', 'Nadam'], loc='lower right')
14 plt.grid(True)
15 plt.savefig('Loss of different optimizer.jpg', dpi=300)
```



In []:

1