



```
In [1]: 1 from keras.datasets import mnist
2 from keras import models
3 from keras import layers
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6 (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

D:\Anaconda3\lib\site-packages\h5py\_init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [2]: 1 train_images = train_images.reshape((60000, 28 * 28))
2 train_images = train_images.astype('float32') / 255
3
4 test_images = test_images.reshape((10000, 28 * 28))
5 test_images = test_images.astype('float32') / 255
6
7 from keras.utils import to_categorical
8
9 train_labels = to_categorical(train_labels)
10 test_labels = to_categorical(test_labels)
```

```
In [3]: 1 def show_train_history(train_history, train, validation, size):
2     plt.plot(train_history.history[train], linewidth=3)
3     plt.plot(train_history.history[validation], linewidth=3)
4     plt.title('Train History')
5     plt.ylabel(train)
6     plt.xlabel('Epoch')
7     plt.xticks([0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19], ['1','2'])
8     plt.legend(['Train', 'Validation'], loc='best')
9     plt.grid(True)
10    if train == 'acc':
11        plt.savefig("batch_size_acc_" + str(size) + ".jpg")
12    if train == 'loss':
13        plt.savefig("batch_size_loss_" + str(size) + ".jpg")
14    plt.show()
```

```
In [4]: 1 #batch_size=64
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 train_history_64 = network.fit(train_images, train_labels, validation_split=0.2,
7                                test_loss, test_acc = network.evaluate(test_images, test_labels)
8                                print('test_loss:', test_loss)
9                                print('test_acc:', test_acc)
```

WARNING:tensorflow:From D:\Anaconda3\lib\site-packages\tensorflow\python\framework\op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From D:\Anaconda3\lib\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 48000 samples, validate on 12000 samples

Epoch 1/20

48000/48000 [=====] - 5s 101us/step - loss: 0.9069 - acc: 0.7986 - val\_loss: 0.4689 - val\_acc: 0.8867

Epoch 2/20

48000/48000 [=====] - 3s 64us/step - loss: 0.4367 - acc: 0.8881 - val\_loss: 0.3633 - val\_acc: 0.9048

Epoch 3/20

48000/48000 [=====] - 3s 64us/step - loss: 0.3646 - acc: 0.9013 - val\_loss: 0.3217 - val\_acc: 0.9127

```

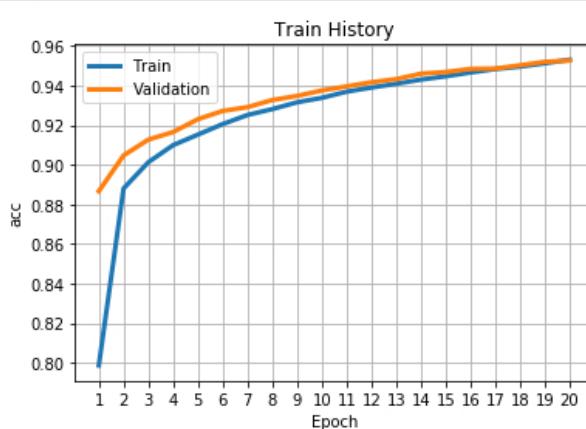
Epoch 4/20
48000/48000 [=====] - 3s 62us/step - loss: 0.3288 - acc: 0.9100 - val_loss: 0.2967 - val_acc: 0.9166
Epoch 5/20
48000/48000 [=====] - 3s 65us/step - loss: 0.3047 - acc: 0.9153 - val_loss: 0.2789 - val_acc: 0.9231
Epoch 6/20
48000/48000 [=====] - 3s 63us/step - loss: 0.2866 - acc: 0.9206 - val_loss: 0.2645 - val_acc: 0.9273
Epoch 7/20
48000/48000 [=====] - 3s 63us/step - loss: 0.2714 - acc: 0.9252 - val_loss: 0.2527 - val_acc: 0.9292
Epoch 8/20
48000/48000 [=====] - 3s 68us/step - loss: 0.2586 - acc: 0.9282 - val_loss: 0.2423 - val_acc: 0.9327
Epoch 9/20
48000/48000 [=====] - 3s 61us/step - loss: 0.2471 - acc: 0.9316 - val_loss: 0.2336 - val_acc: 0.9348
Epoch 10/20
48000/48000 [=====] - 3s 62us/step - loss: 0.2372 - acc: 0.9339 - val_loss: 0.2255 - val_acc: 0.9376
Epoch 11/20
48000/48000 [=====] - 3s 66us/step - loss: 0.2279 - acc: 0.9370 - val_loss: 0.2174 - val_acc: 0.9397
Epoch 12/20
48000/48000 [=====] - 3s 68us/step - loss: 0.2194 - acc: 0.9391 - val_loss: 0.2122 - val_acc: 0.9417
Epoch 13/20
48000/48000 [=====] - 3s 66us/step - loss: 0.2118 - acc: 0.9410 - val_loss: 0.2048 - val_acc: 0.9433
Epoch 14/20
48000/48000 [=====] - 3s 68us/step - loss: 0.2046 - acc: 0.9431 - val_loss: 0.1988 - val_acc: 0.9461
Epoch 15/20
48000/48000 [=====] - 3s 73us/step - loss: 0.1978 - acc: 0.9447 - val_loss: 0.1934 - val_acc: 0.9469
Epoch 16/20
48000/48000 [=====] - 4s 76us/step - loss: 0.1916 - acc: 0.9467 - val_loss: 0.1896 - val_acc: 0.9485
Epoch 17/20
48000/48000 [=====] - 4s 80us/step - loss: 0.1856 - acc: 0.9484 - val_loss: 0.1848 - val_acc: 0.9487
Epoch 18/20
48000/48000 [=====] - 4s 81us/step - loss: 0.1801 - acc: 0.9496 - val_loss: 0.1798 - val_acc: 0.9503
Epoch 19/20
48000/48000 [=====] - 4s 75us/step - loss: 0.1747 - acc: 0.9513 - val_loss: 0.1751 - val_acc: 0.9521
Epoch 20/20
48000/48000 [=====] - 3s 72us/step - loss: 0.1697 - acc: 0.9531 - val_loss: 0.1720 - val_acc: 0.9527
10000/10000 [=====] - 0s 46us/step
test_loss: 0.17141485458612443
test_acc: 0.9505

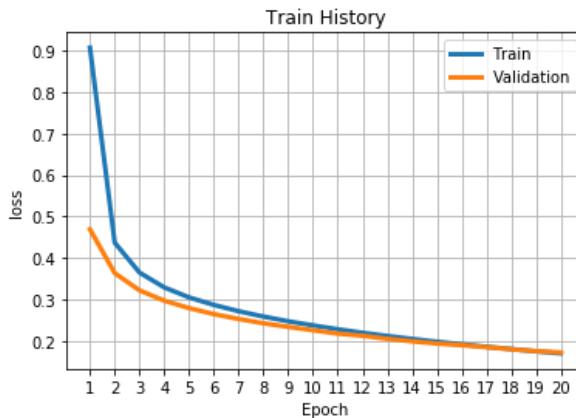
```

```

In [5]: 1 size = 64
2 show_train_history(train_history_64, 'acc', 'val_acc', size)
3 show_train_history(train_history_64, 'loss', 'val_loss', size)
4
5 scores_64 = network.evaluate(train_images, train_labels)
6 print()
7 print("[Info] Accuracy of testing data = {:.2f}%".format(scores_64[1]*100.0))

```





```
60000/60000 [=====] - 3s 50us/step
```

[Info] Accuracy of testing data = 95.3%

```
In [6]: 1 #batch_size=128
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 train_history_128 = network.fit(train_images, train_labels, validation_split=0.2,
7 test_loss, test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', test_loss)
9 print('test_acc:', test_acc)
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/20  
48000/48000 [=====] - 2s 45us/step - loss: 1.2046 - accuracy: 0.7309 - val\_loss: 0.6742 - val\_accuracy: 0.8560

Epoch 2/20  
48000/48000 [=====] - 2s 39us/step - loss: 0.5887 - accuracy: 0.8606 - val\_loss: 0.4705 - val\_accuracy: 0.8849

Epoch 3/20  
48000/48000 [=====] - 2s 38us/step - loss: 0.4642 - accuracy: 0.8811 - val\_loss: 0.3991 - val\_accuracy: 0.8984

Epoch 4/20  
48000/48000 [=====] - 2s 38us/step - loss: 0.4081 - accuracy: 0.8918 - val\_loss: 0.3617 - val\_accuracy: 0.9038

Epoch 5/20  
48000/48000 [=====] - 2s 38us/step - loss: 0.3746 - accuracy: 0.8989 - val\_loss: 0.3370 - val\_accuracy: 0.9101

Epoch 6/20  
48000/48000 [=====] - 2s 39us/step - loss: 0.3511 - accuracy: 0.9035 - val\_loss: 0.3190 - val\_accuracy: 0.9119

Epoch 7/20  
48000/48000 [=====] - 2s 38us/step - loss: 0.3334 - accuracy: 0.9079 - val\_loss: 0.3055 - val\_accuracy: 0.9160

Epoch 8/20  
48000/48000 [=====] - 2s 37us/step - loss: 0.3188 - accuracy: 0.9121 - val\_loss: 0.2937 - val\_accuracy: 0.9193

Epoch 9/20  
48000/48000 [=====] - 2s 39us/step - loss: 0.3068 - accuracy: 0.9148 - val\_loss: 0.2842 - val\_accuracy: 0.9222

Epoch 10/20  
48000/48000 [=====] - 2s 40us/step - loss: 0.2964 - accuracy: 0.9174 - val\_loss: 0.2756 - val\_accuracy: 0.9244

Epoch 11/20  
48000/48000 [=====] - 2s 40us/step - loss: 0.2870 - accuracy: 0.9210 - val\_loss: 0.2679 - val\_accuracy: 0.9256

Epoch 12/20  
48000/48000 [=====] - 2s 39us/step - loss: 0.2784 - accuracy: 0.9233 - val\_loss: 0.2613 - val\_accuracy: 0.9269

Epoch 13/20  
48000/48000 [=====] - 2s 40us/step - loss: 0.2708 - accuracy: 0.9248 - val\_loss: 0.2548 - val\_accuracy: 0.9296

Epoch 14/20  
48000/48000 [=====] - 2s 39us/step - loss: 0.2637 - accuracy: 0.9271 - val\_loss: 0.2490 - val\_accuracy: 0.9303

Epoch 15/20  
48000/48000 [=====] - 2s 38us/step - loss: 0.2570 - accuracy: 0.9286 - val\_loss: 0.2434 - val\_accuracy: 0.9332

Epoch 16/20  
48000/48000 [=====] - 2s 40us/step - loss: 0.2508 - accuracy: 0.9305 - val\_loss: 0.2390 - val\_accuracy: 0.9332

```

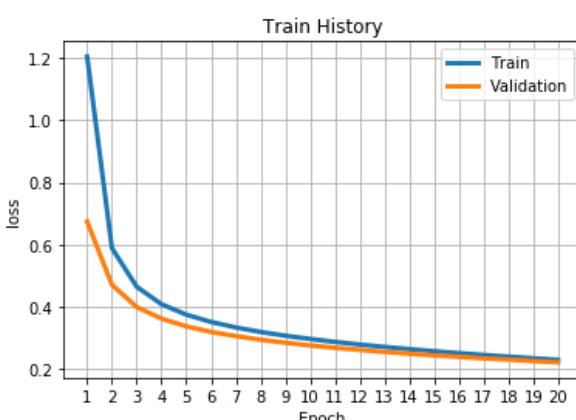
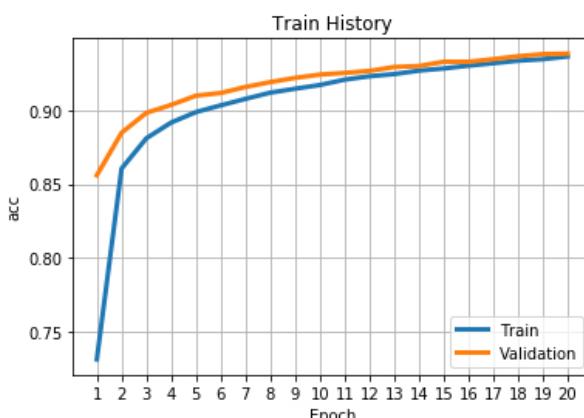
Epoch 17/20
48000/48000 [=====] - 2s 39us/step - loss: 0.2450 - acc: 0.9321 - val_loss: 0.2339 - val_acc: 0.9349
Epoch 18/20
48000/48000 [=====] - 2s 40us/step - loss: 0.2396 - acc: 0.9339 - val_loss: 0.2301 - val_acc: 0.9369
Epoch 19/20
48000/48000 [=====] - 2s 40us/step - loss: 0.2344 - acc: 0.9349 - val_loss: 0.2252 - val_acc: 0.9383
Epoch 20/20
48000/48000 [=====] - 2s 40us/step - loss: 0.2293 - acc: 0.9368 - val_loss: 0.2212 - val_acc: 0.9387
10000/10000 [=====] - 0s 44us/step
test_loss: 0.22186423029005528
test_acc: 0.9392

```

```

In [7]: 1 size = 128
2 show_train_history(train_history_128, 'acc', 'val_acc', size)
3 show_train_history(train_history_128, 'loss', 'val_loss', size)
4
5 scores_128 = network.evaluate(train_images, train_labels)
6 print()
7 print("[Info] Accuracy of testing data = {:.2f}%".format(scores_128[1]*100.0))
8

```



```
60000/60000 [=====] - 2s 41us/step
```

[Info] Accuracy of testing data = 93.8%

```

In [8]: 1 #batch_size=256
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 train_history_256 = network.fit(train_images, train_labels, validation_split=0.2,
7 test_loss, test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', test_loss)
9 print('test_acc:', test_acc)
10

```

```

Train on 48000 samples, validate on 12000 samples
Epoch 1/20
48000/48000 [=====] - 1s 27us/step - loss: 1.6042 - acc: 0.6296 - val_loss: 1.0446 - val_acc: 0.8163
Epoch 2/20
48000/48000 [=====] - 1s 22us/step - loss: 0.8617 - acc: 0.8244 - val_loss: 0.6791 - val_acc: 0.8591

```

```

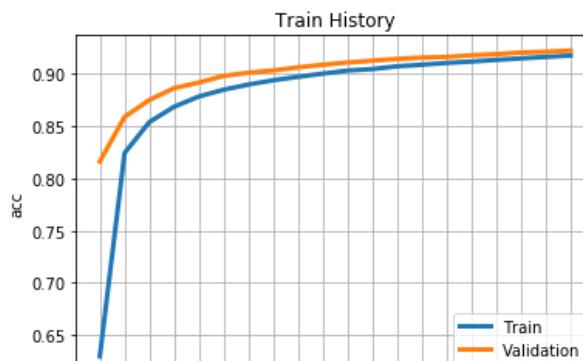
Epoch 3/20
48000/48000 [=====] - 1s 22us/step - loss: 0.6387 - ac
c: 0.8541 - val_loss: 0.5425 - val_acc: 0.8752
Epoch 4/20
48000/48000 [=====] - 1s 22us/step - loss: 0.5393 - ac
c: 0.8688 - val_loss: 0.4728 - val_acc: 0.8866
Epoch 5/20
48000/48000 [=====] - 1s 22us/step - loss: 0.4823 - ac
c: 0.8786 - val_loss: 0.4301 - val_acc: 0.8920
Epoch 6/20
48000/48000 [=====] - 1s 24us/step - loss: 0.4448 - ac
c: 0.8851 - val_loss: 0.4003 - val_acc: 0.8983
Epoch 7/20
48000/48000 [=====] - 1s 22us/step - loss: 0.4178 - ac
c: 0.8901 - val_loss: 0.3794 - val_acc: 0.9014
Epoch 8/20
48000/48000 [=====] - 1s 23us/step - loss: 0.3971 - ac
c: 0.8941 - val_loss: 0.3626 - val_acc: 0.9035
Epoch 9/20
48000/48000 [=====] - 1s 23us/step - loss: 0.3809 - ac
c: 0.8974 - val_loss: 0.3492 - val_acc: 0.9065
Epoch 10/20
48000/48000 [=====] - 1s 22us/step - loss: 0.3673 - ac
c: 0.9005 - val_loss: 0.3386 - val_acc: 0.9090
Epoch 11/20
48000/48000 [=====] - 1s 22us/step - loss: 0.3559 - ac
c: 0.9034 - val_loss: 0.3290 - val_acc: 0.9112
Epoch 12/20
48000/48000 [=====] - 1s 24us/step - loss: 0.3459 - ac
c: 0.9049 - val_loss: 0.3210 - val_acc: 0.9129
Epoch 13/20
48000/48000 [=====] - 1s 23us/step - loss: 0.3372 - ac
c: 0.9075 - val_loss: 0.3137 - val_acc: 0.9146
Epoch 14/20
48000/48000 [=====] - 1s 22us/step - loss: 0.3294 - ac
c: 0.9090 - val_loss: 0.3072 - val_acc: 0.9158
Epoch 15/20
48000/48000 [=====] - 1s 23us/step - loss: 0.3224 - ac
c: 0.9107 - val_loss: 0.3019 - val_acc: 0.9165
Epoch 16/20
48000/48000 [=====] - 1s 22us/step - loss: 0.3160 - ac
c: 0.9121 - val_loss: 0.2962 - val_acc: 0.9181
Epoch 17/20
48000/48000 [=====] - 1s 23us/step - loss: 0.3101 - ac
c: 0.9136 - val_loss: 0.2912 - val_acc: 0.9191
Epoch 18/20
48000/48000 [=====] - 1s 22us/step - loss: 0.3046 - ac
c: 0.9151 - val_loss: 0.2867 - val_acc: 0.9208
Epoch 19/20
48000/48000 [=====] - 1s 22us/step - loss: 0.2995 - ac
c: 0.9165 - val_loss: 0.2825 - val_acc: 0.9214
Epoch 20/20
48000/48000 [=====] - 1s 22us/step - loss: 0.2946 - ac
c: 0.9179 - val_loss: 0.2786 - val_acc: 0.9226
10000/10000 [=====] - 0s 46us/step
test_loss: 0.27871255747675894
test_acc: 0.9238

```

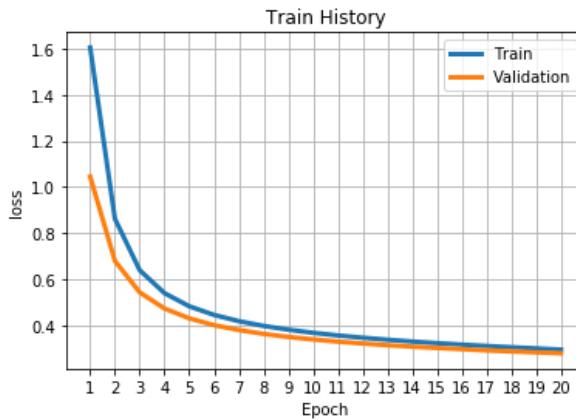
```

In [9]: 1 size = 256
2 show_train_history(train_history_256, 'acc', 'val_acc', size)
3 show_train_history(train_history_256, 'loss', 'val_loss', size)
4
5 scores_256 = network.evaluate(train_images, train_labels)
6 print()
7 print("[Info] Accuracy of testing data = {:.2f}%".format(scores_256[1]*100.0))

```



```
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20  
Epoch
```



```
60000/60000 [=====] - 3s 45us/step
```

```
[Info] Accuracy of testing data = 92.0%
```

```
In [10]: 1 #batch_size=512
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['ac
6 train_history_512 = network.fit(train_images, train_labels, validation_split=
7 test_loss, test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', test_loss)
9 print('test_acc:', test_acc)
```

```
Train on 48000 samples, validate on 12000 samples
Epoch 1/20
48000/48000 [=====] - 1s 23us/step - loss: 1.8946 - ac
c: 0.5212 - val_loss: 1.5104 - val_acc: 0.7456
Epoch 2/20
48000/48000 [=====] - 1s 17us/step - loss: 1.2810 - ac
c: 0.7745 - val_loss: 1.0482 - val_acc: 0.8213
Epoch 3/20
48000/48000 [=====] - 1s 17us/step - loss: 0.9487 - ac
c: 0.8199 - val_loss: 0.8087 - val_acc: 0.8504
Epoch 4/20
48000/48000 [=====] - 1s 17us/step - loss: 0.7721 - ac
c: 0.8421 - val_loss: 0.6760 - val_acc: 0.8617
Epoch 5/20
48000/48000 [=====] - 1s 17us/step - loss: 0.6680 - ac
c: 0.8550 - val_loss: 0.5936 - val_acc: 0.8735
Epoch 6/20
48000/48000 [=====] - 1s 17us/step - loss: 0.6003 - ac
c: 0.8638 - val_loss: 0.5385 - val_acc: 0.8797
Epoch 7/20
48000/48000 [=====] - 1s 17us/step - loss: 0.5526 - ac
c: 0.8701 - val_loss: 0.4991 - val_acc: 0.8853
Epoch 8/20
48000/48000 [=====] - 1s 17us/step - loss: 0.5171 - ac
c: 0.8755 - val_loss: 0.4692 - val_acc: 0.8893
Epoch 9/20
48000/48000 [=====] - 1s 18us/step - loss: 0.4896 - ac
c: 0.8796 - val_loss: 0.4458 - val_acc: 0.8930
Epoch 10/20
48000/48000 [=====] - 1s 17us/step - loss: 0.4676 - ac
c: 0.8833 - val_loss: 0.4271 - val_acc: 0.8958
Epoch 11/20
48000/48000 [=====] - 1s 17us/step - loss: 0.4494 - ac
c: 0.8862 - val_loss: 0.4115 - val_acc: 0.8985
Epoch 12/20
48000/48000 [=====] - 1s 17us/step - loss: 0.4342 - ac
c: 0.8884 - val_loss: 0.3985 - val_acc: 0.9003
Epoch 13/20
48000/48000 [=====] - 1s 17us/step - loss: 0.4211 - ac
c: 0.8906 - val_loss: 0.3874 - val_acc: 0.9023
Epoch 14/20
48000/48000 [=====] - 1s 19us/step - loss: 0.4098 - ac
c: 0.8924 - val_loss: 0.3778 - val_acc: 0.9038
Epoch 15/20
48000/48000 [=====] - 1s 18us/step - loss: 0.3997 - ac
c: 0.8944 - val_loss: 0.3691 - val_acc: 0.9054
```

```

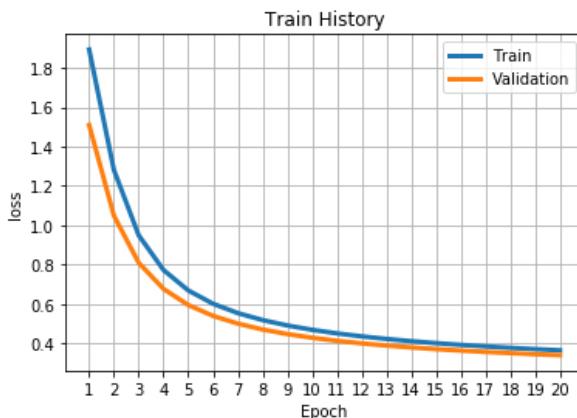
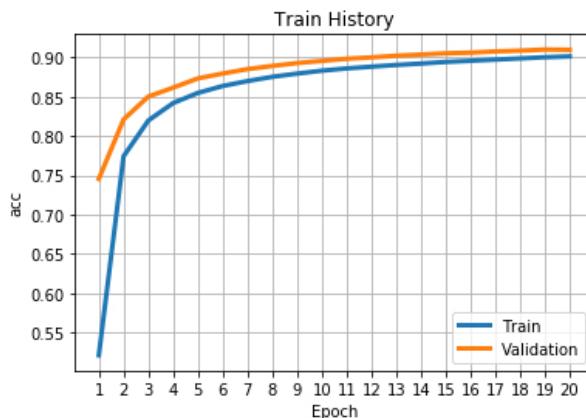
Epoch 16/20
48000/48000 [=====] - 1s 18us/step - loss: 0.3908 - acc: 0.8959 - val_loss: 0.3619 - val_acc: 0.9062
Epoch 17/20
48000/48000 [=====] - 1s 17us/step - loss: 0.3830 - acc: 0.8974 - val_loss: 0.3549 - val_acc: 0.9078
Epoch 18/20
48000/48000 [=====] - 1s 17us/step - loss: 0.3758 - acc: 0.8989 - val_loss: 0.3491 - val_acc: 0.9088
Epoch 19/20
48000/48000 [=====] - 1s 17us/step - loss: 0.3692 - acc: 0.9004 - val_loss: 0.3432 - val_acc: 0.9102
Epoch 20/20
48000/48000 [=====] - 1s 17us/step - loss: 0.3632 - acc: 0.9015 - val_loss: 0.3382 - val_acc: 0.9099
10000/10000 [=====] - 0s 49us/step
test_loss: 0.3410462572097778
test_acc: 0.9079

```

```

In [11]: 1 size = 512
2 show_train_history(train_history_512, 'acc', 'val_acc', size)
3 show_train_history(train_history_512, 'loss', 'val_loss', size)
4
5 scores_512 = network.evaluate(train_images, train_labels)
6 print()
7 print("[Info] Accuracy of testing data = {:.2f}%".format(scores_512[1]*100.0))
8

```



```
60000/60000 [=====] - 3s 47us/step
```

[Info] Accuracy of testing data = 90.4%

```

In [12]: 1 #batch_size=1024
2 network = models.Sequential()
3 network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
4 network.add(layers.Dense(10, activation='softmax'))
5 network.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
6 train_history_1024 = network.fit(train_images, train_labels, validation_split=0.2)
7 test_loss, test_acc = network.evaluate(test_images, test_labels)
8 print('test_loss:', test_loss)
9 print('test_acc:', test_acc)
10

```

```

Train on 48000 samples, validate on 12000 samples
Epoch 1/20
48000/48000 [=====] - 1s 20us/step - loss: 2.0580 - acc: 0.3625 - val_loss: 1.8132 - val_acc: 0.6155

```

```

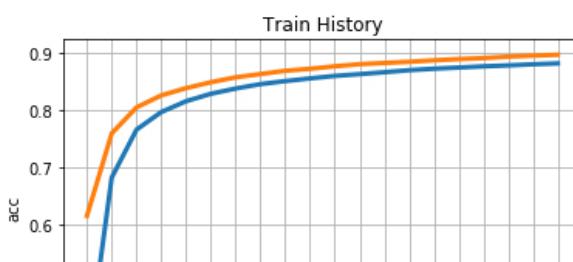
Epoch 2/20
48000/48000 [=====] - 1s 13us/step - loss: 1.6529 - acc: 0.6825 - val_loss: 1.4608 - val_acc: 0.7597
Epoch 3/20
48000/48000 [=====] - 1s 13us/step - loss: 1.3537 - acc: 0.7663 - val_loss: 1.2007 - val_acc: 0.8052
Epoch 4/20
48000/48000 [=====] - 1s 13us/step - loss: 1.1373 - acc: 0.7970 - val_loss: 1.0154 - val_acc: 0.8262
Epoch 5/20
48000/48000 [=====] - 1s 13us/step - loss: 0.9833 - acc: 0.8160 - val_loss: 0.8837 - val_acc: 0.8388
Epoch 6/20
48000/48000 [=====] - 1s 13us/step - loss: 0.8725 - acc: 0.8290 - val_loss: 0.7877 - val_acc: 0.8492
Epoch 7/20
48000/48000 [=====] - 1s 13us/step - loss: 0.7906 - acc: 0.8382 - val_loss: 0.7163 - val_acc: 0.8578
Epoch 8/20
48000/48000 [=====] - 1s 13us/step - loss: 0.7282 - acc: 0.8459 - val_loss: 0.6615 - val_acc: 0.8635
Epoch 9/20
48000/48000 [=====] - 1s 13us/step - loss: 0.6793 - acc: 0.8511 - val_loss: 0.6182 - val_acc: 0.8692
Epoch 10/20
48000/48000 [=====] - 1s 13us/step - loss: 0.6400 - acc: 0.8559 - val_loss: 0.5832 - val_acc: 0.8727
Epoch 11/20
48000/48000 [=====] - 1s 14us/step - loss: 0.6077 - acc: 0.8602 - val_loss: 0.5543 - val_acc: 0.8771
Epoch 12/20
48000/48000 [=====] - 1s 15us/step - loss: 0.5808 - acc: 0.8633 - val_loss: 0.5303 - val_acc: 0.8808
Epoch 13/20
48000/48000 [=====] - 1s 14us/step - loss: 0.5579 - acc: 0.8666 - val_loss: 0.5097 - val_acc: 0.8829
Epoch 14/20
48000/48000 [=====] - 1s 14us/step - loss: 0.5382 - acc: 0.8702 - val_loss: 0.4921 - val_acc: 0.8849
Epoch 15/20
48000/48000 [=====] - 1s 13us/step - loss: 0.5210 - acc: 0.8728 - val_loss: 0.4767 - val_acc: 0.8875
Epoch 16/20
48000/48000 [=====] - 1s 14us/step - loss: 0.5059 - acc: 0.8749 - val_loss: 0.4633 - val_acc: 0.8897
Epoch 17/20
48000/48000 [=====] - 1s 13us/step - loss: 0.4925 - acc: 0.8770 - val_loss: 0.4514 - val_acc: 0.8914
Epoch 18/20
48000/48000 [=====] - 1s 13us/step - loss: 0.4806 - acc: 0.8786 - val_loss: 0.4406 - val_acc: 0.8940
Epoch 19/20
48000/48000 [=====] - 1s 13us/step - loss: 0.4698 - acc: 0.8805 - val_loss: 0.4310 - val_acc: 0.8958
Epoch 20/20
48000/48000 [=====] - 1s 13us/step - loss: 0.4600 - acc: 0.8821 - val_loss: 0.4224 - val_acc: 0.8973
10000/10000 [=====] - 0s 46us/step
test_loss: 0.4272878079652786
test_acc: 0.8952

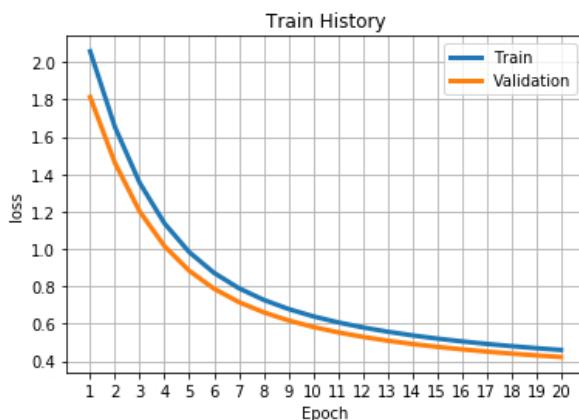
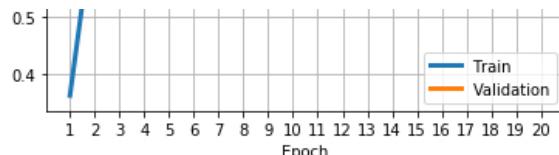
```

```

In [13]: 1 size = 1024
2 show_train_history(train_history_1024, 'acc', 'val_acc', size)
3 show_train_history(train_history_1024, 'loss', 'val_loss', size)
4
5
6 scores_1024 = network.evaluate(train_images, train_labels)
7 print()
8 print("[Info] Accuracy of testing data = {:.2f}%".format(scores_1024[1]*100))

```

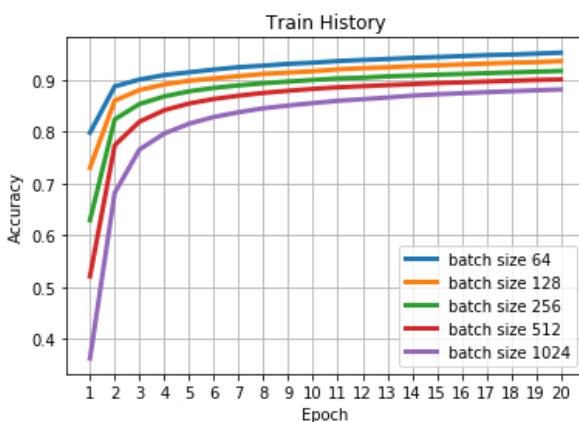




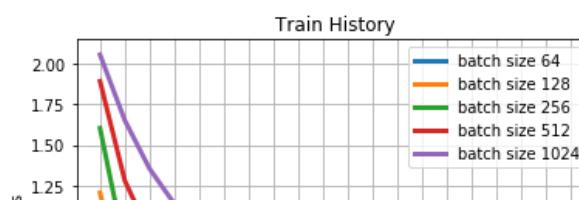
```
60000/60000 [=====] - 3s 45us/step
```

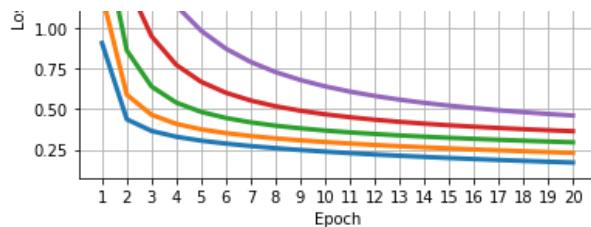
[Info] Accuracy of testing data = 88.6%

```
In [14]: 1 plt.plot(train_history_64.history['acc'], linewidth=3)
2 plt.plot(train_history_128.history['acc'], linewidth=3)
3 plt.plot(train_history_256.history['acc'], linewidth=3)
4 plt.plot(train_history_512.history['acc'], linewidth=3)
5 plt.plot(train_history_1024.history['acc'], linewidth=3)
6 plt.title('Train History')
7 plt.xlabel('Epoch')
8 plt.ylabel('Accuracy')
9 plt.xticks([0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19], ['1','2','3'])
10 plt.legend(['batch size 64','batch size 128','batch size 256','batch size 512','batch size 1024'])
11 plt.grid(True)
12 plt.savefig('batch_size_acc_all.jpg',dpi=300)
```

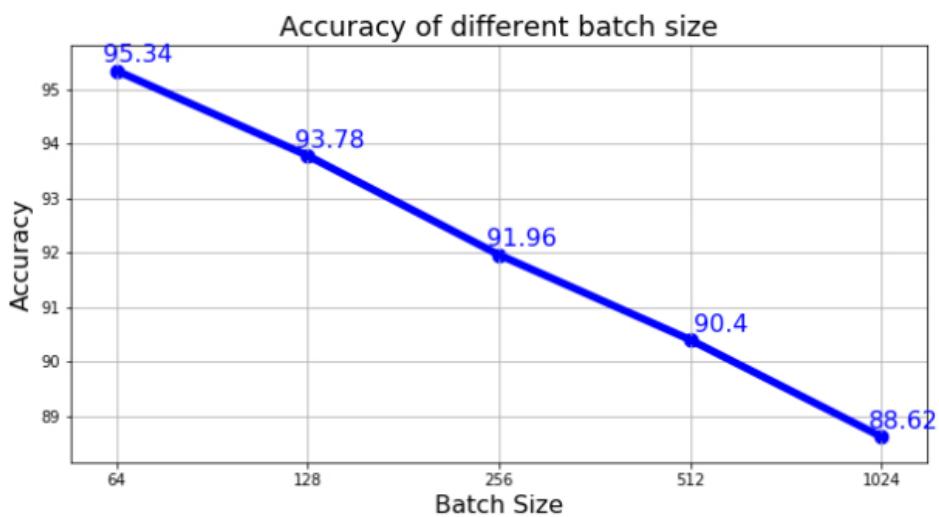


```
In [15]: 1 plt.plot(train_history_64.history['loss'], linewidth=3)
2 plt.plot(train_history_128.history['loss'], linewidth=3)
3 plt.plot(train_history_256.history['loss'], linewidth=3)
4 plt.plot(train_history_512.history['loss'], linewidth=3)
5 plt.plot(train_history_1024.history['loss'], linewidth=3)
6 plt.title('Train History')
7 plt.xlabel('Epoch')
8 plt.ylabel('Loss')
9 plt.xticks([0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19], ['1','2','3'])
10 plt.legend(['batch size 64','batch size 128','batch size 256','batch size 512','batch size 1024'])
11 plt.grid(True)
12 plt.savefig('batch_size_loss_all.jpg',dpi=300)
```





```
In [53]: 1 acc = [round(scores_64[1]*100,2),round(scores_128[1]*100,2),round(scores_256[1]*100,2),round(scores_512[1]*100,2),round(scores_1024[1]*100,2)]
2 x = ['64','128','256','512','1024']
3 plt.figure(figsize=(10,5))
4 plt.plot(x,acc,'b',lw=5)
5 plt.scatter(x, acc, s = 75,color='b')
6 plt.title('Accuracy of different batch size', fontsize='18')
7 plt.xlabel('Batch Size',fontsize=16)
8 plt.ylabel('Accuracy',fontsize=16)
9 plt.grid(True)
10 for x,y in enumerate(acc):
11     plt.text(x+0.3,y+0.15,'%s %y, ha='right', color='b',fontsize=16)
12 plt.savefig('Accuracy of different batch size.jpg',dpi=300)
```



```
In [ ]: 1
```