

InceptionV3

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 from keras.applications.inception_v3 import InceptionV3
4 from keras.applications.inception_v3 import preprocess_input
5 from keras.preprocessing.image import ImageDataGenerator, array_to_img
6 from keras.preprocessing import image
7 from keras.models import Model, Sequential
8 from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D, AveragePooling2D
9 from keras.layers import Activation, Dropout, Flatten, Dense
10 from keras.callbacks import ModelCheckpoint, EarlyStopping
11 from keras.optimizers import SGD
12 from keras import backend as K
```

D:\Anaconda\lib\site-packages\h5py_init_.py:36: FutureWarning: Conversion of the second argument of issubtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
from _conv import register_converters as _register_converters
Using TensorFlow backend.

```
In [2]: 1 # 影像縮放
2 img_width, img_height = 299, 299
3
4 train_data_dir = 'D:/Anaconda3/Scripts/5 上課資料/電腦視覺與人機互動/20191031HW/cats_and_dogs_small/train'
5 validation_data_dir = 'D:/Anaconda3/Scripts/5 上課資料/電腦視覺與人機互動/20191031HW/cats_and_dogs_small/validation'
6 nb_train_samples = 2000
7 nb_validation_samples = 800
8 batch_size = 2
```

```
In [3]: 1 if K.image_data_format() == 'channels_first':
2     input_shape = (3, img_width, img_height)
3 else:
4     input_shape = (img_width, img_height, 3)
```

```
In [4]: 1 base_model = InceptionV3(weights='imagenet', include_top=False)
2
3 # 添加全連接層平均池化層
4 x = base_model.output
5 x = GlobalAveragePooling2D()(x)
6 # 增加全連接層
7 x = Dense(1024, activation='relu', name='fc1')(x)
8 prediction = Dense(2, activation='softmax', name='predictions')(x)
9 model = Model(inputs=base_model.input, outputs=prediction)
10 model.summary()
```

concatenate_2 (Concatenate)	(None, None, None, 7 0)	activation_92[0][0] activation_93[0][0]
activation_94 (Activation)	(None, None, None, 1 0)	batch_normalization_94[0][0]
mixed10 (Concatenate)	(None, None, None, 2 0)	activation_86[0][0] mixed9_1[0][0] concatenate_2[0][0] activation_94[0][0]
global_average_pooling2d_1 (Glo (None, 2048)	0	mixed10[0][0]
fc1 (Dense)	(None, 1024)	2098176 global_average_pooling2d_1[0][0]
predictions (Dense)	(None, 2)	2050 fc1[0][0]

Total params: 23,903,610
Trainable params: 23,868,578
Non-trainable params: 34,432

```
In [5]: 1 # 矩結所有層(除了Bottleneck Layers之外)來進行微調
2 for layer in model.layers:
3     if layer.name in ['predictions']:
4         continue
5     layer.trainable = False
```

```
In [6]: 1 df = pd.DataFrame([layer.name, layer.trainable] for layer in model.layers), columns=['layer', 'trainable'])
```

```
Out[6]:
```

	layer	trainable
0	input_1	False
1	conv2d_1	False
2	batch_normalization_1	False
3	activation_1	False
4	conv2d_2	False
5	batch_normalization_2	False
6	activation_2	False
7	conv2d_3	False
8	batch_normalization_3	False
9	activation_3	False
10	max_pooling2d_1	False
11	conv2d_4	False

```
In [7]: 1 train_datagen = ImageDataGenerator(rotation_range=40,
2 width_shift_range=0.2,
3 height_shift_range=0.2,
4 shear_range=0.2,
5 zoom_range=0.2,
6 horizontal_flip=True,
7 fill_mode='nearest')
8 train_generator = train_datagen.flow_from_directory(directory='D:/Anaconda3/Scripts/5 上課資料/電腦視覺與人機互動/20191031HW/cats_and_dogs_small/train',
9 target_size=(img_width, img_height),
10 batch_size=batch_size,
11 class_mode='categorical')
12
13 validation_datagen = ImageDataGenerator()
14 validation_generator = validation_datagen.flow_from_directory(directory='D:/Anaconda3/Scripts/5 上課資料/電腦視覺與人機互動/20191031HW/cats_and_dogs_small/validation',
15 target_size=(img_width, img_height),
16 batch_size=batch_size,
17 class_mode='categorical')
```

Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.

```
In [8]: 1 model.compile(optimizer = 'rmsprop', loss = 'categorical_crossentropy', metrics = ['accuracy'])
2 history = model.fit_generator(
3     train_generator,
4     steps_per_epoch = train_generator.samples // batch_size,
5     validation_data = validation_generator,
6     validation_steps = validation_generator.samples // batch_size,
7     epochs = 20)
8
9 model.save('model-InceptionV3-final1.h5')
```

```
1000/1000 [=====] - 131ms/step - loss: 0.4315 - acc: 0.7765 - val_loss: 7.5340 - val_acc: 0.518
0
Epoch 15/20
1000/1000 [=====] - 132s 132ms/step - loss: 0.4087 - acc: 0.8070 - val_loss: 7.5696 - val_acc: 0.515
00.4
Epoch 16/20
1000/1000 [=====] - 132s 132ms/step - loss: 0.4517 - acc: 0.7640 - val_loss: 7.7938 - val_acc: 0.511
0
Epoch 17/20
1000/1000 [=====] - 132s 132ms/step - loss: 0.4343 - acc: 0.7815 - val_loss: 7.7535 - val_acc: 0.511
0
Epoch 18/20
```

```
1000/1000 [=====] - 131s 131ms/step - loss: 0.4261 - acc: 0.7895 - val_loss: 7.9017 - val_acc: 0.506
0
Epoch 19/20
1000/1000 [=====] - 131s 131ms/step - loss: 0.4490 - acc: 0.7710 - val_loss: 7.8034 - val_acc: 0.512
0
Epoch 20/20
1000/1000 [=====] - 131s 131ms/step - loss: 0.4296 - acc: 0.7815 - val_loss: 7.7969 - val_acc: 0.512
0
```

```
In [9]: 1 for i, layer in enumerate(base_model.layers):
2     print(i, layer.name)
3 for layer in model.layers[172:]:
4     layer.trainable = False
5 for layer in model.layers[172:]:
6     layer.trainable = True
```

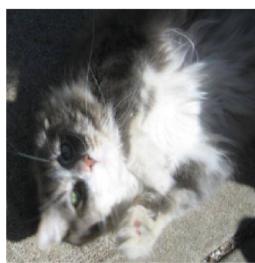
```
0 input_1
1 conv2d_1
2 batch_normalization_1
3 activation_1
4 conv2d_2
5 batch_normalization_2
6 activation_2
7 conv2d_3
8 batch_normalization_3
9 activation_3
10 max_pooling2d_1
11 conv2d_4
12 batch_normalization_4
13 activation_4
14 conv2d_5
15 batch_normalization_5
16 activation_5
17 max_pooling2d_2
18 conv2d_9
19 batch_normalization_9
```

```
In [10]: 1 from keras.optimizers import SGD
2 model.compile(optimizer=SGD(lr=0.001, momentum=0.9), loss='categorical_crossentropy', metrics=['accuracy'])
3 history = model.fit_generator(
4     train_generator,
5     steps_per_epoch = train_generator.samples // batch_size,
6     validation_data = validation_generator,
7     validation_steps = validation_generator.samples // batch_size,
8     epochs = 20)
9
10 model.save('model-inceptionv3-final12.h5')
```

```
Epoch 1/20
1000/1000 [=====] - 239s 239ms/step - loss: 1.2781 - acc: 0.6990 - val_loss: 7.8385 - val_acc: 0.5100
Epoch 2/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.3086 - acc: 0.8605 - val_loss: 8.0590 - val_acc: 0.5000
Epoch 3/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.1861 - acc: 0.9325 - val_loss: 7.8990 - val_acc: 0.5020
Epoch 4/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.1582 - acc: 0.9355 - val_loss: 6.4260 - val_acc: 0.5210
Epoch 5/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.1392 - acc: 0.9425 - val_loss: 8.0293 - val_acc: 0.5010
Epoch 6/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.1198 - acc: 0.9550 - val_loss: 7.4270 - val_acc: 0.5150
Epoch 7/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.1033 - acc: 0.9585 - val_loss: 8.0533 - val_acc: 0.5000
Epoch 8/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.1081 - acc: 0.9590 - val_loss: 8.0567 - val_acc: 0.5000
Epoch 9/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.0858 - acc: 0.9675 - val_loss: 8.0590 - val_acc: 0.5000
Epoch 10/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.0875 - acc: 0.9670 - val_loss: 8.0590 - val_acc: 0.5000
Epoch 11/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.0691 - acc: 0.9730 - val_loss: 8.0590 - val_acc: 0.5000
Epoch 12/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0698 - acc: 0.9705 - val_loss: 8.0659 - val_acc: 0.4990
Epoch 13/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0762 - acc: 0.9720 - val_loss: 7.6983 - val_acc: 0.4990
Epoch 14/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0469 - acc: 0.9805 - val_loss: 7.4622 - val_acc: 0.5140
Epoch 15/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0530 - acc: 0.9805 - val_loss: 8.0576 - val_acc: 0.5000
Epoch 16/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0465 - acc: 0.9810 - val_loss: 6.4573 - val_acc: 0.5450
Epoch 17/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0504 - acc: 0.9790 - val_loss: 6.4857 - val_acc: 0.5570
Epoch 18/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.0435 - acc: 0.9815 - val_loss: 6.9573 - val_acc: 0.5370
Epoch 19/20
1000/1000 [=====] - 232s 232ms/step - loss: 0.0325 - acc: 0.9875 - val_loss: 7.5819 - val_acc: 0.5080
Epoch 20/20
1000/1000 [=====] - 231s 231ms/step - loss: 0.0431 - acc: 0.9845 - val_loss: 8.0590 - val_acc: 0.5000
```

```
In [11]: 1 from IPython.display import display
2 import matplotlib.pyplot as plt
3
4 X_val_sample, _ = next(validation_generator)
5 y_pred = model.predict(X_val_sample)
6
7 nb_sample = 4
8 for X, y in zip(X_val_sample[:nb_sample], y_pred[:nb_sample]):
9     s = pd.Series({'Cat': np.max(y), 'Dog': np.max(y)})
10    axes = s.plot(kind='bar')
11    axes.set_xlabel('Class')
12    axes.set_ylabel('Probability')
13    axes.set_ylim([0, 1])
14    plt.show()
15
16    img = array_to_img(X)
17    display(img)
```

<Figure size 640x480 with 1 Axes>



<Figure size 640x480 with 1 Axes>

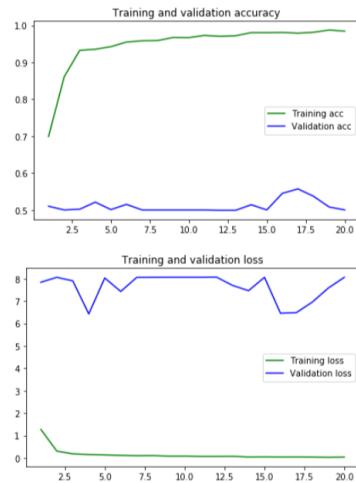


```
In [12]: 1 scores = model.evaluate_generator(generator=validation_generator, steps=validation_generator.samples // batch_size)
2 print("\nLoss: %.2f, Accuracy: %.2f%%" % (scores[0]/100, scores[1]*100))
```

Loss: 0.08, Accuracy: 50.00%

```
In [13]: 1 modelname = 'Inceptionv3'
2
3 acc = history.history['acc']
4 val_acc = history.history['val_acc']
5 loss = history.history['loss']
6 val_loss = history.history['val_loss']
7
8 epochs = range(1, len(acc) + 1)
9
10 metrics = 'acc'
11 filename = 'image/' + modelname + '_' + metrics
12 plt.plot(epochs, acc, 'g', label='Training acc')
13 plt.plot(epochs, val_acc, 'b', label='Validation acc')
14 plt.title('Training and validation accuracy')
15 plt.legend(loc='best')
16 plt.tight_layout()
17 plt.savefig(filename, dpi=300)
18 print('Save',filename)
19 plt.figure()
20
21 metrics = 'loss'
22 filename = 'image/' + modelname + '_' + metrics
23 plt.plot(epochs, loss, 'g', label='Training loss')
24 plt.plot(epochs, val_loss, 'b', label='Validation loss')
25 plt.title('Training and validation loss')
26 plt.legend(loc='best')
27 plt.tight_layout()
28 plt.savefig(filename, dpi=300)
29 print('Save',filename)
30 plt.show()
```

Save image/Inceptionv3_acc
Save image/Inceptionv3_loss



In []:

1